

```
In [3]: import tensorflow as tf
        from tensorflow import keras
        from keras.models import Sequential
        from keras.layers import Activation, Dense, Flatten, BatchNormalization, Co
```

```
In [4]: model = Sequential()

        model.add(Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_s
        model.add(MaxPool2D(pool_size=(2, 2), strides=2))

        model.add(Conv2D(filters=64, kernel_size=(3, 3), activation='relu', padding
        model.add(MaxPool2D(pool_size=(2, 2), strides=2))

        model.add(Conv2D(filters=128, kernel_size=(3, 3), activation='relu', paddin
        model.add(MaxPool2D(pool_size=(2, 2), strides=2))

        model.add(Flatten())

        model.add(Dense(64,activation = "relu"))
        model.add(Dense(128,activation = "relu"))
        #model.add(Dropout(0.2))
        model.add(Dense(128,activation = "relu"))
        #model.add(Dropout(0.3))
        model.add(Dense(10,activation = "softmax"))
```

```
In [5]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_1 (Conv2D)	(None, 31, 31, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 64)	0
conv2d_2 (Conv2D)	(None, 13, 13, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 128)	0
flatten (Flatten)	(None, 4608)	0
dense (Dense)	(None, 64)	294976
dense_1 (Dense)	(None, 128)	8320
dense_2 (Dense)	(None, 128)	16512
dense_3 (Dense)	(None, 10)	1290
=====		
Total params: 414346 (1.58 MB)		
Trainable params: 414346 (1.58 MB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [6]: model.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy',
                    reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=1, min_delta=0.001),
                    early_stop = EarlyStopping(monitor='val_loss', min_delta=0, patience=2, verbose=1))

model.compile(optimizer=SGD(learning_rate=0.001), loss='categorical_crossentropy',
                    reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=1, min_delta=0.001),
                    early_stop = EarlyStopping(monitor='val_loss', min_delta=0, patience=2, verbose=1))
```

```
In [10]: history2 = model.fit(train_batches, epochs=10, callbacks=[reduce_lr, early_
```

Epoch 1/10

252/252 [=====] - 47s 180ms/step - loss: 2.8262e-04 - accuracy: 1.0000 - val_loss: 3.7159e-04 - val_accuracy: 1.0000 - lr: 5.0000e-04

Epoch 2/10

252/252 [=====] - 38s 151ms/step - loss: 2.6475e-04 - accuracy: 1.0000 - val_loss: 3.4914e-04 - val_accuracy: 1.0000 - lr: 5.0000e-04

Epoch 3/10

252/252 [=====] - 33s 129ms/step - loss: 2.4882e-04 - accuracy: 1.0000 - val_loss: 3.2906e-04 - val_accuracy: 1.0000 - lr: 5.0000e-04

Epoch 4/10

252/252 [=====] - 32s 128ms/step - loss: 2.3455e-04 - accuracy: 1.0000 - val_loss: 3.1113e-04 - val_accuracy: 1.0000 - lr: 5.0000e-04

Epoch 5/10

252/252 [=====] - 32s 128ms/step - loss: 2.2174e-04 - accuracy: 1.0000 - val_loss: 2.9502e-04 - val_accuracy: 1.0000 - lr: 5.0000e-04

Epoch 6/10

252/252 [=====] - 33s 132ms/step - loss: 2.1018e-04 - accuracy: 1.0000 - val_loss: 2.8033e-04 - val_accuracy: 1.0000 - lr: 5.0000e-04

Epoch 7/10

252/252 [=====] - 33s 131ms/step - loss: 1.9971e-04 - accuracy: 1.0000 - val_loss: 2.6704e-04 - val_accuracy: 1.0000 - lr: 5.0000e-04

Epoch 8/10

252/252 [=====] - 775s 3s/step - loss: 1.9018e-04 - accuracy: 1.0000 - val_loss: 2.5491e-04 - val_accuracy: 1.0000 - lr: 5.0000e-04

Epoch 9/10

252/252 [=====] - 37s 145ms/step - loss: 1.8148e-04 - accuracy: 1.0000 - val_loss: 2.4376e-04 - val_accuracy: 1.0000 - lr: 5.0000e-04

Epoch 10/10

252/252 [=====] - 35s 140ms/step - loss: 1.7351e-04 - accuracy: 1.0000 - val_loss: 2.3352e-04 - val_accuracy: 1.0000 - lr: 5.0000e-04


```

In [9]: import tensorflow as tf
from tensorflow import keras
from keras.models import Sequential
from keras.layers import Activation, Dense, Flatten, BatchNormalization, Conv2D
from keras.optimizers import Adam, SGD
from keras.metrics import categorical_crossentropy
from keras.preprocessing.image import ImageDataGenerator
import itertools
import random
import warnings
import numpy as np
import cv2
import matplotlib.pyplot as plt
from keras.callbacks import ReduceLROnPlateau
from keras.callbacks import ModelCheckpoint, EarlyStopping
warnings.simplefilter(action='ignore', category=FutureWarning)

train_path = r'D:\gesture\train'
test_path = r'D:\gesture\test'

train_batches = ImageDataGenerator(preprocessing_function=tf.keras.applications.vgg16.preprocess_image)
test_batches = ImageDataGenerator(preprocessing_function=tf.keras.applications.vgg16.preprocess_image)

imgs, labels = next(train_batches)

#Plotting the images...
def plotImages(images_arr):
    fig, axes = plt.subplots(1, 10, figsize=(30,20))
    axes = axes.flatten()
    for img, ax in zip( images_arr, axes):
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
        ax.imshow(img)
        ax.axis('off')
    plt.tight_layout()
    plt.show()

plotImages(imgs)
print(imgs.shape)
print(labels)

model = Sequential()

model.add(Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 3)))
model.add(MaxPool2D(pool_size=(2, 2), strides=2))

model.add(Conv2D(filters=64, kernel_size=(3, 3), activation='relu', padding='same'))
model.add(MaxPool2D(pool_size=(2, 2), strides=2))

model.add(Conv2D(filters=128, kernel_size=(3, 3), activation='relu', padding='same'))
model.add(MaxPool2D(pool_size=(2, 2), strides=2))

model.add(Flatten())

model.add(Dense(64,activation = "relu"))
model.add(Dense(128,activation = "relu"))
#model.add(Dropout(0.2))
model.add(Dense(128,activation = "relu"))
#model.add(Dropout(0.3))

```

```

model.add(Dense(10,activation ="softmax"))

# In[23]:

model.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy',
               reduce_lr = ReduceLRonPlateau(monitor='val_loss', factor=0.2, patience=1, min_delta=0.001),
               early_stop = EarlyStopping(monitor='val_loss', min_delta=0, patience=2, verbose=0))

model.compile(optimizer=SGD(learning_rate=0.001), loss='categorical_crossentropy',
               reduce_lr = ReduceLRonPlateau(monitor='val_loss', factor=0.2, patience=1, min_delta=0.001),
               early_stop = EarlyStopping(monitor='val_loss', min_delta=0, patience=2, verbose=0))

history2 = model.fit(train_batches, epochs=10, callbacks=[reduce_lr, early_stop],
                     imgs, labels = next(train_batches) # For getting next batch of imgs...

imgs, labels = next(test_batches) # For getting next batch of imgs...
scores = model.evaluate(imgs, labels, verbose=0)
print(f'{model.metrics_names[0]} of {scores[0]}; {model.metrics_names[1]} of {scores[1]}')

#model.save('best_model_dataflair.h5')
model.save('best_model_dataflair3.h5')

print(history2.history)

imgs, labels = next(test_batches)

model = keras.models.load_model(r"best_model_dataflair3.h5")

scores = model.evaluate(imgs, labels, verbose=0)
print(f'{model.metrics_names[0]} of {scores[0]}; {model.metrics_names[1]} of {scores[1]}')

model.summary()

scores #[loss, accuracy] on test data...
model.metrics_names

word_dict = {0:'One',1:'Ten',2:'Two',3:'Three',4:'Four',5:'Five',6:'Six',7:'Seven'}

predictions = model.predict(imgs, verbose=0)
print("predictions on a small set of test data--")
print("")
for ind, i in enumerate(predictions):
    print(word_dict[np.argmax(i)], end=' ')

plotImages(imgs)
print('Actual labels')
for i in labels:
    print(word_dict[np.argmax(i)], end=' ')

print(imgs.shape)

history2.history

```

Found 2520 images belonging to 10 classes.
Found 490 images belonging to 10 classes.

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
In [15]: # For getting next batch of testing imgs...
         imgs, labels = next(test_batches)

         scores = model.evaluate(imgs, labels, verbose=0)
         print(f'{model.metrics_names[0]} of {scores[0]}; {model.metrics_names[1]} o

         #Once the model is fitted we save the model using model.save() function.

         model.save(r'C:\Users\vishwanth\Downloads\sign-language-recognition-project

loss of 0.00024940239381976426; accuracy of 100.0%
```

```
In [13]: word_dict = {0:'One',1:'Ten',2:'Two',3:'Three',4:'Four',5:'Five',6:'Six',7:
        predictions = model.predict(imgs, verbose=0)
        print("predictions on a small set of test data--")
        print("")
        for ind, i in enumerate(predictions):
            print(word_dict[np.argmax(i)], end='  ')

        plotImages(imgs)
        print('Actual labels')
        for i in labels:
            print(word_dict[np.argmax(i)], end='  ')
```

predictions on a small set of test data--

Eight Ten Ten Seven Three Ten Six One Ten Seven

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



Actual labels

Eight Ten Ten Seven Three Ten Six One Ten Seven


```

In [21]: import numpy as np
import cv2
import keras
from keras.preprocessing.image import ImageDataGenerator
import tensorflow as tf

model = keras.models.load_model(r'C:\Users\vishwanth\Downloads\sign-languag

background = None
accumulated_weight = 0.5

ROI_top = 100
ROI_bottom = 300
ROI_right = 150
ROI_left = 350

def cal_accum_avg(frame, accumulated_weight):

    global background

    if background is None:
        background = frame.copy().astype("float")
        return None

    cv2.accumulateWeighted(frame, background, accumulated_weight)

def segment_hand(frame, threshold=25):
    global background

    diff = cv2.absdiff(background.astype("uint8"), frame)

    _, thresholded = cv2.threshold(diff, threshold, 255, cv2.THRESH_BINARY)

    #Fetching contours in the frame (These contours can be of hand or any o
    contours, hierarchy = cv2.findContours(thresholded.copy(), cv2.RETR_EXT

    # If Length of contours list = 0, means we didn't get any contours...
    if len(contours) == 0:
        return None
    else:
        # The Largest external contour should be the hand
        hand_segment_max_cont = max(contours, key=cv2.contourArea)

        # Returning the hand segment(max contour) and the thresholded image
        return (thresholded, hand_segment_max_cont)

cam = cv2.VideoCapture(0)
num_frames = 0
while True:
    ret, frame = cam.read()

    # flipping the frame to prevent inverted image of captured frame...
    frame = cv2.flip(frame, 1)

```

```

frame_copy = frame.copy()

# ROI from the frame
roi = frame[ROI_top:ROI_bottom, ROI_right:ROI_left]

gray_frame = cv2.cvtColor(roi, cv2.COLOR_BGR2GRAY)
gray_frame = cv2.GaussianBlur(gray_frame, (9, 9), 0)

if num_frames < 70:

    cal_accum_avg(gray_frame, accumulated_weight)

    cv2.putText(frame_copy, "FETCHING BACKGROUND...PLEASE WAIT", (80, 4

else:
    # segmenting the hand region
    hand = segment_hand(gray_frame)

    # Checking if we are able to detect the hand...
    if hand is not None:

        thresholded, hand_segment = hand

        # Drawing contours around hand segment
        cv2.drawContours(frame_copy, [hand_segment + (ROI_right, ROI_to

        cv2.imshow("Thesholded Hand Image", thresholded)

        thresholded = cv2.resize(thresholded, (64, 64))
        thresholded = cv2.cvtColor(thresholded, cv2.COLOR_GRAY2RGB)
        thresholded = np.reshape(thresholded, (1,thresholded.shape[0],t

        pred = model.predict(thresholded)
        cv2.putText(frame_copy, word_dict[np.argmax(pred)], (170, 45),

    # Draw ROI on frame_copy
    cv2.rectangle(frame_copy, (ROI_left, ROI_top), (ROI_right, ROI_bottom),

    # incrementing the number of frames for tracking
    num_frames += 1

    # Display the frame with segmented hand
    cv2.putText(frame_copy, "DataFlair hand sign recognition_ _ _", (10, 20
    cv2.imshow("Sign Detection", frame_copy)

    # Close windows with Esc
    k = cv2.waitKey(1) & 0xFF

    if k == 27:
        break

# Release the camera and destroy all the windows
cam.release()
cv2.destroyAllWindows()

```

```
1/1 [=====] - 0s 228ms/step
1/1 [=====] - 0s 58ms/step
1/1 [=====] - 0s 61ms/step
1/1 [=====] - 0s 73ms/step
1/1 [=====] - 0s 54ms/step
1/1 [=====] - 0s 61ms/step
1/1 [=====] - 0s 53ms/step
1/1 [=====] - 0s 61ms/step
1/1 [=====] - 0s 67ms/step
1/1 [=====] - 0s 61ms/step
1/1 [=====] - 0s 57ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 73ms/step
1/1 [=====] - 0s 62ms/step
1/1 [=====] - 0s 57ms/step
1/1 [=====] - 0s 60ms/step
1/1 [=====] - 0s 66ms/step
1/1 [=====] - 0s 57ms/step
1/1 [=====] - 0s 63ms/step
```

In []:

In []: